

# Performance Evaluation

Final project defence

---

Timon Zimmermann & Jimi Vaubien

February 13, 2019

École Polytechnique Fédéral de Lausanne

Evaluate and compare performances of different **in-memory** databases systems

# Overall schema

Maintain all the data efficiently in memory ...

Order	Country	Product	Sales
456	France	corn	1000
457	Italy	wheat	900
458	Spain	rice	600
459	Italy	rice	800
460	Denmark	corn	500
461	Denmark	rice	600
462	Belgium	rice	600
463	Italy	rice	1100



... and keep the cores busy



Figure 1: How it works

## A few points about such systems

- ▶ *In theory*, relies on the main-memory to entirely store the data
- ▶ *In practice*, often used in pair with on-disk databases
  - Parts with critical response-time on main-memory
  - Non-critical parts on slower databases (HDD, SSD)
- ▶ Thread not idled by CPU because of the speed
- ▶ Uses persistent databases as back-up

Compare the atomic operations “set/get” of 2 different in-memory database systems.

# Databases systems compared

- ▶ Memcached (<http://memcached.org/>)
  - Written in: C
  - Used by Google, Facebook...
- ▶ Redis (<https://redis.io/>)
  - Written in: C
  - Used by Twitter, GitHub...

We use 3 different metrics:

- ▶ Response time (`get` & `set` operations)
- ▶ CPU usage
- ▶ RAM usage

Also investigate effects on the `system`:

- ▶ Similarly to Homework 3
- ▶ Congestion collapse

We use 3 different metrics:

- ▶ Response time (`get` & `set` operations)
- ▶ `CPU` usage
- ▶ ~~RAM usage~~

Also investigate effects on the `system`:

- ▶ Similarly to Homework 3
- ▶ ~~Congestion collapse~~



# Where?

- ▶ AWS (*Amazon*) EC2 instance
  - Burstable CPU compromising results
  - Limited bandwidth at reasonable pricing
- ▶ Own cluster connected through *1 Gb/s Ethernet*
  - “Attacker” with 6 physical cores (12 logical)
  - Database system hosted on 2 physical cores (4 logical) and 1 Gb RAM
  - Isolated in a virtual machine to minimize interaction with other processes

# Operations performed

- ▶ Get and Set operations
- ▶ Scrambled and Contiguous
- ▶ 200 clients, each 200 requests/second
  - Less: Too easy to handle
  - More: Bigger load than “attacker” CPUs permit
- ▶ Operations distributed following  $\text{Zipf}(0, 100'000, 0.99)$  [1]

# Results: Redis end-to-end

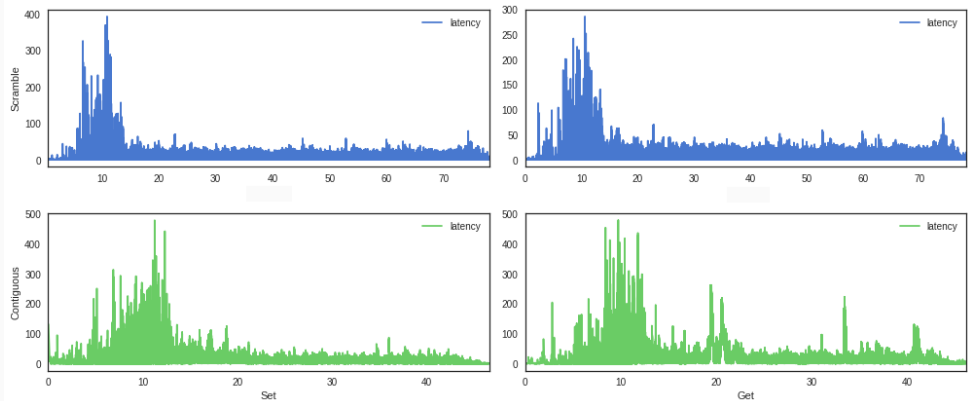


Figure 2: End to end latency time serie




Figure 3: End to end latency distribution




Figure 4: End to end latency box plots

# Results: Redis stable regime




Figure 5: Stable regime latency distribution




Figure 6: Stable regime latency box plots

# Results: Memcached end-to-end

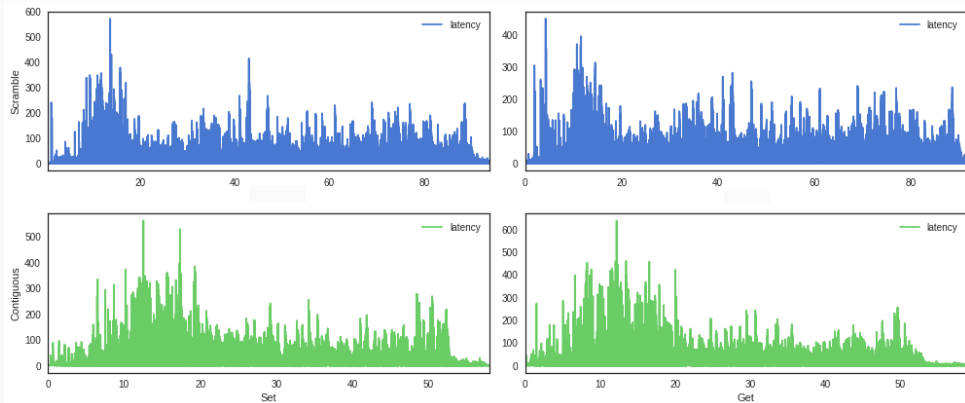


Figure 7: End to end latency time serie






Figure 8: End to end latency distribution




Figure 9: End to end latency box plots

# Results: Memcached stable regime




Figure 10: Stable regime latency distribution




Figure 11: Stable regime latency box plots

# Comparisons: Impact on the system

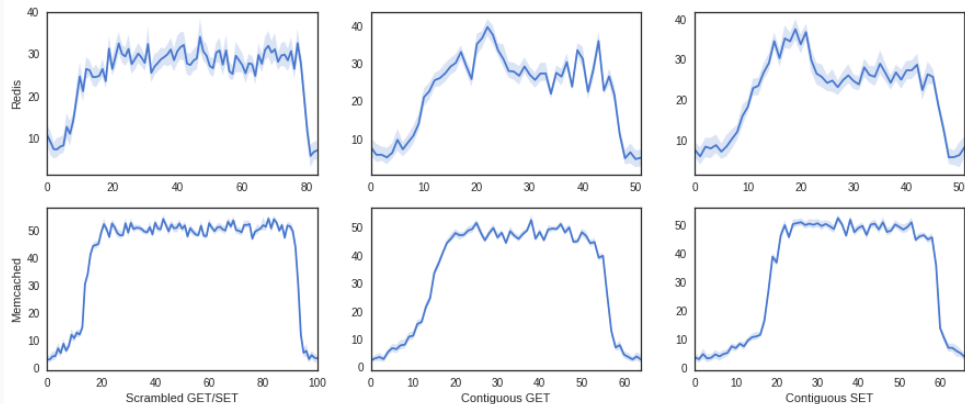


Figure 12: CPU usage of Redis and Memcached for different operations





Figure 13: RAM usage of Redis and Memcached for different operations


- ▶ Redis uses more RAM than Memcached
- ▶ Memcached CPU usage is higher but stays almost the same once in stationary regime
  - Easier to predict power consumption for huge data centers

# Comparisons: Latency end-to-end





# Comparisons: Latency in “stable” regime



- ▶ Redis  $\approx 1.8x$  faster than Memcached in 2 categories
  - One category even (contiguous SET)
  - Memcached  $\approx 2.2x$  faster than Redis for **contiguous GET**
  - Memcached widely used for very-high **get** load (Facebook, Google and so on)
- ▶ Redis **optimizations** out-of-the-box ?
  - Using more RAM to pre-load instruction set
  - Default optimizations (RAM hash tables and so on), where Memcached is **“lower-level”**
- ▶ Confidence intervals validating the “instruction pre-loading” hypothesis
  - Redis latency less “sparse”

# Possible improvements and extensions

- ▶ Test with optimizations disabled and fully enabled on both
- ▶ More cluster power to reach system's limit
  - Current setup bottlenecked by Network interface controller
- ▶ Bigger payloads to inspect how RAM is arranged and optimized

Questions?



E. F. S. J. B. Atikoglu, Y. Xu and M. Paleczny, “Workload analysis of a large-scale key-value store,” in *Proceedings of the SIGMETRICS’12*, June 2012.